

UNIVERSITAT DE BARCELONA

FUNDAMENTALS OF DATA SCIENCE MASTER'S THESIS

Man-made Structures Detection from Space

Author:

Eduard RIBAS FERNÁNDEZ

Supervisor:

Dr. Jordi VITRIÀ

*A thesis submitted in partial fulfillment of the requirements
for the degree of MSc in Fundamentals of Data Science*

in the

Facultat de Matemàtiques i Informàtica

June 20, 2019

UNIVERSITAT DE BARCELONA

Abstract

Facultat de Matemàtiques i Informàtica

Man-made Structures Detection from Space

by Eduard RIBAS FERNÁNDEZ

The Thesis Abstract is written here (and usually kept to just this page). The page is kept centered vertically so can expand into the blank space above the title too...

Acknowledgements

The acknowledgments and the people to thank go here, don't forget to include your project advisor...

Contents

Abstract	iii
Acknowledgements	v
1 Introduction	3
1.1 Motivation	3
1.2 Satellogic	3
1.3 Thesis Outline	3
2 Deep Learning	5
2.1 Introduction to Deep Learning	5
2.2 Convolutional Neural Networks	7
2.3 CNN Architectures	7
3 Building Datasets	9
3.1 Requirements and Considerations	9
3.2 Existing Datasets	9
3.3 Google Maps	10
3.4 USGS, Land Cover	11
3.4.1 Getting the Data	11
3.4.2 Data Processing and Labeling	13
4 Deep Learning Approach	17
4.1 Image features	17
4.1.1 Traditional image features	17
4.1.2 NN features and transfer learning	17
4.1.3 pretrained resnet, weights on imagenet	18
4.2 Proposed architecture	18
4.2.1 resnet activations	18
4.2.2 complete architecture	18
4.2.3 training	18
5 Results	19
5.1 Results	19
5.1.1 Subsection 1	19
5.2 Cost and Environment impact	19
6 Conclusions	21
6.1 Conclusions	21
6.1.1 Subsection 1	21
6.2 Further work	21

7	Introduction	23
7.0.1	Tables	23
7.0.2	Figures	23
7.0.3	Typesetting mathematics	24
7.1	Sectioning and Subsectioning	25
	Bibliography	27

Chapter 1

Introduction

1.1 Motivation

BigEarthNet Motivation: Existing Remote Sensing Datasets contain a small number of annotated images and therefore do not suffice to train a complex ConvNet with many different parameters. Remote sensing versus Computer (Using pretrained nets on ImageNet)

BagOfVisualWords: useless

Extreme value theory-based calibration: useless

Climate Change, population control, illegal habitat/illegal land use

1.2 Satellogic

1.3 Thesis Outline

Chapter 2

Deep Learning

In this chapter, we will provide a short overview of the theoretical concepts and recent advances in Deep Learning. We will give a basic introduction to neural networks, where we will discuss convolutional neural networks in more detail as these are the type of algorithms utilized in this work. We further will summarize some of the most popular convolutional neural network architectures.

2.1 Introduction to Deep Learning

Deep Learning (DL) models have led to vast performance improvements in a large variety of domains, and therefore have gained substantial popularity over the last decade. These models were initially inspired by the human brain and analogies in neuroscience, which is why this class of algorithms is called neural networks (NN). The two most popular neural network architectures are convolutional neural networks (CNN) and recurrent neural networks (RNN). CNNs have driven major breakthroughs in visual object recognition [12], and image [9], video [17] and audio [11] processing while RNNs brought about advances in research on sequential data, i.e. in speech and text [4]. However, the superior performance of Neural Networks compared to traditional machine learning algorithms is not limited to the aforementioned domains. Other fields in which NNs have advanced the state-of-the-art, for instance, include bioinformatics [14] and the analysis of data from elementary particle physics [3].

CITE IN THIS SECTION Neural networks define a class of models that are composed of a variable number of processing layers (Hidden units) of simple models, and are generally used to map a fixed size input (e.g. the pixels of an image) to a fixed size output (e.g. a category or a probability). A Hidden unit of a fully connected neural network has connections between all the nodes of the previous layer and the next layer. These connections are fully parametrized by the weights of the network. In analogy to a firing neuron, a non-linear activation function is applied to the output of every Hidden unit. Historically, sigmoid and tanh have been used as activation functions. Nowadays the most popular activation function is the rectified linear unit (ReLU) $f(x) = \max(0, x)$ [7]. We will see the reason for this at the end of the section. In Fig. 2.1 we show an example of a fully connected 3-layer neural network.

The strength of neural networks is that they can learn arbitrarily complex non-linear input-output mappings [5], and that they have the ability to automatically extract features from raw data. The latter is in stark contrast to traditional machine learning algorithms, which require careful feature engineering. For instance, when dealing with images, the multi layer architecture of neural networks allows to learn

different features at every stage of the network, where the complexity of the features increases as the network gets deeper.

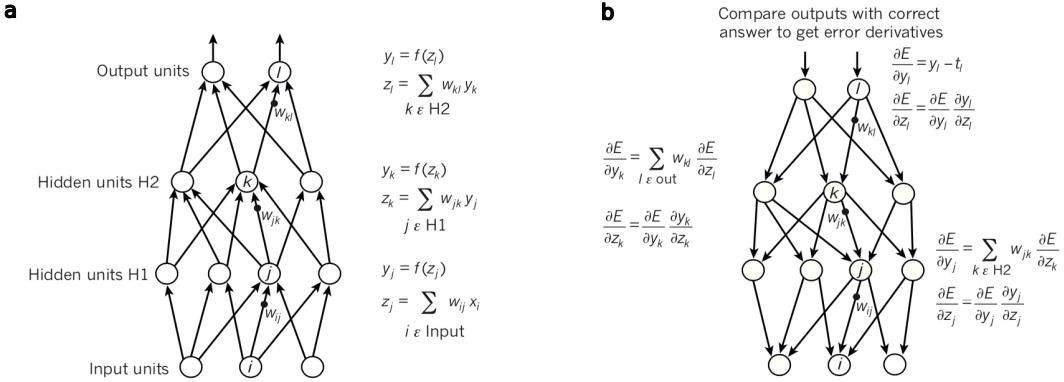


FIGURE 2.1: Example of 3-Layer Neural Network. (a) Feed-forward representation of a neural network with two hidden units H1 and H2 and a binary output unit. The inputs to every layer are weighted averages, specified by the weights w , of the outputs of the previous layer. In every layer, the outputs are generated by applying a non-linear function to the inputs. The most popular function for this purpose is the ReLu (see text). **(b)** Back-propagation of the error in order to learn the optimal weights of the neural network. The error is quantified by a loss function E at the output of the neural network, which is a measure of the discrepancy between the desired output and the actual output of the network. During backpropagation the chain-rule is applied recursively to the loss function in a backward manner. Specifically, at every layer the derivative of the error with respect to the inputs is computed by multiplying the upstream gradient with the local gradient. The upstream gradient is the derivative of the loss function with respect to the output of each unit, which is a weighted sum of the input derivatives of the layer above. The local gradient is the derivative of the non-linear function $f(z)$ with respect to its inputs. Starting from the output of the network one finally obtains the derivative of the loss function with respect to all weights, so that the network can minimize the loss by adjusting the weights. Panel is adapted from [13].

ic gradient descent [2] or a variant of it. The computation of the gradient of the objective function with respect to all weights of the network is accomplished using backpropagation [15] (see Fig. 2.1(b)). Intuitively, INTUITIVE EXPLANATION ABOUT BACKPROPAGATION.

Althoug artifical neural networks have been known and studied since the 1950s, the first real breakthroughs were accomplished after the year 2009. The main reason for the specticism towards learning complex multilayer networks stemmed form the believe, that it was unfeasible to learn the parameters. However, in the last decade several many factors lead up to the explosion of deep learning and applications of which we will name some. First, with the advent of Big Data there suddenly there was enough data available to learn the hundreds of millions of parameters of a typical deep neural network. Second, the rectified linear unit solved the vanishing gradient problem, so that much deeper networks could be trained. And last but not least, powerful graphics processing units (GPU) have turned out to be the ideal hardware to train large neural networks. All of these advances have led to the development of new network architectures, amongst one of them the convolutional neural networks.

2.2 Convolutional Neural Networks

2.3 CNN Architectures

Chapter 3

Building Datasets

In this chapter, we will give an overview of existing (labelled) aerial imagery datasets and outline the reasons why none of them is suitable for our investigation. Following this discussion, we will describe two approaches for obtaining our own labelled dataset.

3.1 Requirements and Considerations

Before we go into the presentation of existing labelled datasets we discuss the requirements that the dataset needs to fulfill in order to serve for the investigation in this thesis project. As a refresher, we want to detect human impact on aerial images and determine the dependency on resolution per pixel of a chosen evaluation metric. Ideally, the range for the resolutions should scale from a few tens of centimeters to a few tens of meters, whereas the images with low resolution can be generated from the high resolution images by downsampling. Having in mind previous arguments, we mainly need to consider three aspects.

First, we need to have imagery data with labels that can be used to clearly distinguish between existing and non-existing human impact, respectively. This impact might be classified pixel wise, or as binary classification for the entire image, or as multi-class classification that can be translated into binary labelling. Second, we need a balanced dataset of approximately the same number of images for both labels, and variations of the images as large as possible with respect to different terrains. Third, the images need to have a resolution per pixel which is equal or better than 1m. Also, the height and width of the images should measure at least 500×500 pixels, so that one has enough room for downsampling.

3.2 Existing Datasets

In table 3.1 we have summarized the most relevant remote sensing datasets with ground truth labels, that can be found in literature. The table lists the name of the dataset together with the bibliographic reference. It also details the data source for the images. Further it contains a description about the number of images, the resolution of the images, the size (in pixel) of the images where images are squared, and the number of categories.

The datasets were collected using different publicly available data sources. These range from pure low resolution satellite imagery (Sentinel-2) to high-resolution images taken with an aircraft (USGS) to a mix of different image sources (Google Earth).

The satellite images have a resolution of equal or larger than 10 m and they are collected with the Sentinel-2 satellites of the European Earth observation program Copernicus. Although the datasets from this source (BigEarthNet and EuroSat) are

comparatively large, they do not suffice for our purpose, because the resolution is not good enough and the images are too small.

name	source	images	resolution (m)	size (pixel)	categories
BigEarthNet [16]	Sentinel-2	590,326	10, 20, 60	120, 60, 20	~ 50
EuroSAT [10]	Sentinel-2	27,000	10	64	10
UCMerced [19]	USGS	2100	0.3	256	21
DeepSat [1]	USGS	405,000	1	28	6
AID [18]	Google Earth	10,000	0.5 - 8	600	30
PatternNet [20]	Google Earth	30,400	0.06 - 4.69	256	38

TABLE 3.1: Publicly available remote sensing datasets with labels.

The USGS National Map Urban Area Imagery collection ([see link](#)) was utilized to collect remote sensing datasets in the two works UCMerced and DeepSat, where the former is the dataset that comes closest to our requirements. It has 21 categories of which only 2 belong to images without human impact, while the other 19 show human impact. The DeepSat dataset unfortunately consists of image patches which are only 28×28 large, so that we aren't able to study these images as a function of resolution.

The datasets using Google Earth as data source are collected using either the Google Earth or the Google Maps application programming interface (API). These images vary in resolution as well as in their original data provider since Google accesses several data sources. Both datasets, the AID and the PatternNet dataset, have about 30 categories with several images in each category. Here, different categories have different resolutions per pixel, and again most of the categories relate to urban areas so that we do not have sufficient images without human impact. Even the categories that in principle should not show human influence contain images that break this rule.

Overall, the main issue with these datasets stems from the fact that none of them was collected with the purpose to analyze the human footprint and therefore they are very unbalanced, and do not contain sufficient variety of images for the classes without human influence. Therefore, we decided to collect and label images by ourselves. In our first approach we used the Google Maps API, and in our final approach we used datasets from the USGS Aerial Imagery collection.

3.3 Google Maps

Google has a public API that allows for querying images from their service Google Maps. In its most basic form, the API accepts as input parameters a latitude (lat) and longitude (lon), a zoom, and the image size (in pixels) to return. Given this set of parameters one can calculate the resolution in meter per pixel [8], which is given by

$$\frac{\text{meter}}{\text{pixel}} = \frac{156543.03392 \cdot \cos(\frac{\text{latitude} \cdot \pi}{180})}{2^{\text{zoom}}} \quad (3.1)$$

Equipped with this toolkit, we developed an automated image download pipeline that was based on one of several approaches of selecting and downloading images in a given area. In our first approach we selected images that were Gaussian distributed around a center location defined by a set of lat/lon coordinates. We further provided a precision parameter (standard deviation of the Gaussian), the number

of images to download, a list of zooms, the desired image size, and the number of pixels to crop from the border of the image in order to remove e.g. the Google logo. Another approach consisted in downloading randomly sampled location from within a rectangle defined by its upper left and lower right lat/long coordinates, respectively.

Although any of these two approaches would have served to build a complete dataset in an automated fashion, we finally decided to use a different data source due to the following reasons. According to our advisor from Satellogic, Google Maps images have one major drawback regarding the resolution per pixel. The visible Google Maps image is an interpolation from different spectral bands, where the gray band has a higher resolution than the RGB bands. Therefore, the resolution estimated by Eq. 3.1 is not representing the actual pixel resolution for the three color bands. We did not further investigate into this issue, and instead turned to a different solution, which is discussed next.

3.4 USGS, Land Cover

3.4.1 Getting the Data

To be able to construct a balanced and representative dataset we first decided to focus on images of the United States, which allows for a large variety of different terrains. We then used as data source the Aerial Imagery datasets from USGS Earth-explorer ([see link](#)) which we combined with information about Land Cover and Land Use available from the USGS Land Cover Viewer ([see link](#)).



FIGURE 3.1: Example images of category Agriculture. All images in this figure show clear signs of human impact. The images have a size of 512×512 pixels and a resolution of 0.3m per pixel.

When looking for images we excluded cities and highly developed urban areas, and instead focussed on unpopulated areas. Specifically, we limited our image search to the four Land Use categories Agriculture, Shrubland-Grassland, Semi-Desert, Forest-Woodland that can be found in the USGS Land Cover Viewer. Note that these categories served as a rough geographic reference to pin down geolocations of interest, in order to guarantee a dataset with a good variety of different

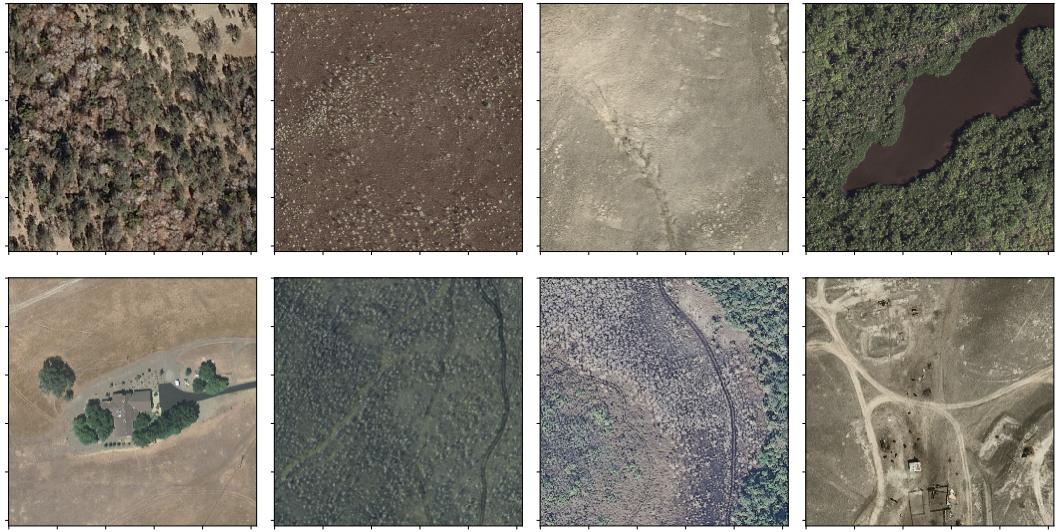


FIGURE 3.2: Example images of category Shrubland-grassland. The images in the first row do not contain any human influence, while the images in the second row show influence by humans. The images in this figure have a size of 512×512 pixels and a resolution of 0.3m per pixel.

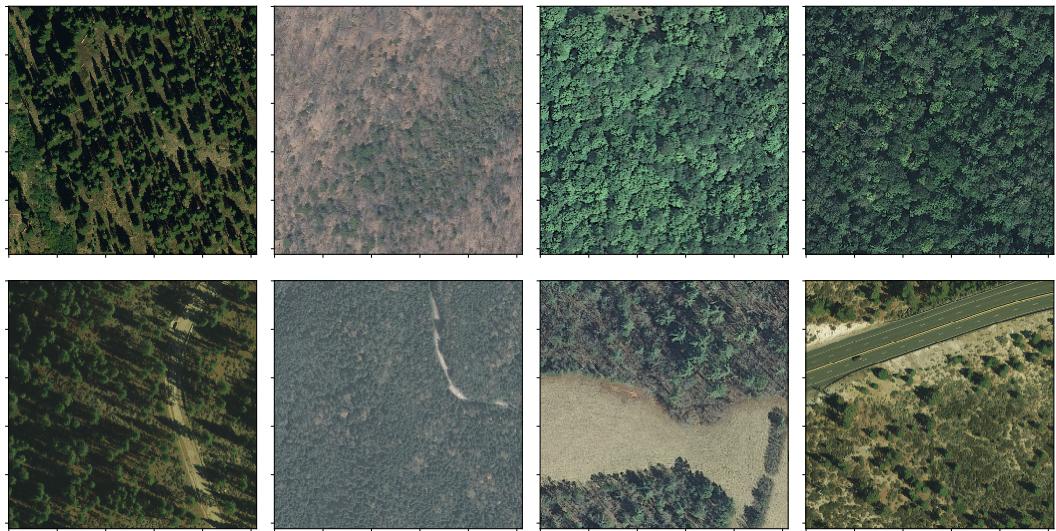


FIGURE 3.3: Example images of category Forest-woodland. The images in the first row do not contain any human influence, while the images in the second row show influence by humans. The images in this figure have a size of 512×512 pixels and a resolution of 0.3m per pixel.

terrains. We also found that it is harder to find images without human, which is why we selected many images from national parks. However, within a given area/terrain we always tried to have images with and without impact.

Once an area was pointed out as a region of interest, we located it on USGS Earth-explorer and downloaded images from that area. In particular, we constructed two datasets with 0.3m and 1m resolution, respectively. The former was taken from the category High Resolution Orthoimagery and the latter from the category National

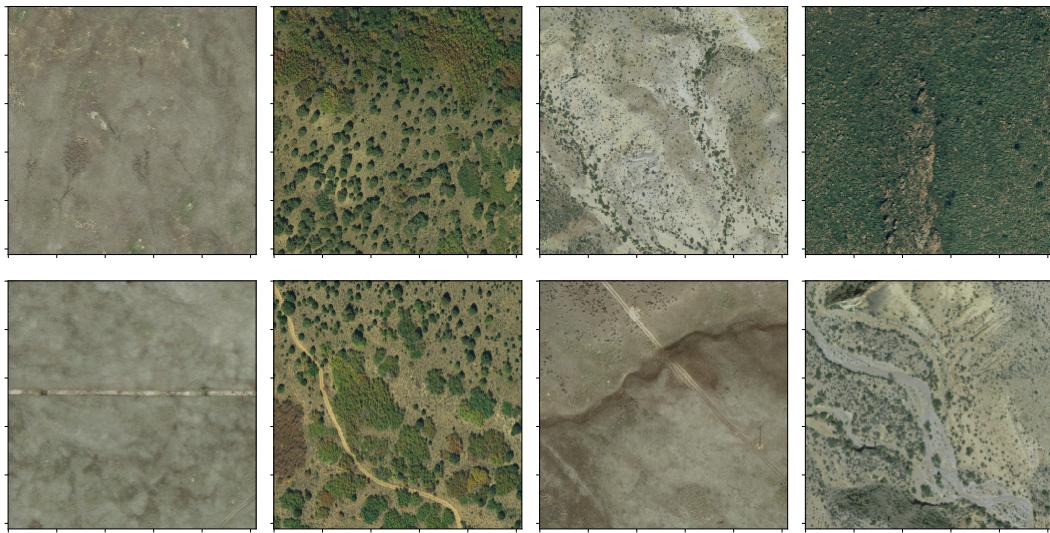


FIGURE 3.4: Example images of category Semi-desert. The images in the first row do not contain any human influence, while the images in the second row show influence by humans. The images in this figure have a size of 512×512 pixels and a resolution of 0.3m per pixel.

Agriculture Imagery Program (NAIP). Note that the images in these categories usually have a height and a width of several thousand pixels, and hence occupy a few hundreds of Megabytes of disk space. We cropped smaller images out of the raw images, which will be discussed in more detail in the following section. Overall, we downloaded about 100 images for each dataset.

3.4.2 Data Processing and Labeling

Our data processing pipeline consists of the following steps:

- Download large raw images
- Crop images of size 512×512 pixels
- Label images with either zero (no human impact), one (minimal human impact), two (obvious human impact)
- Degrade images, i.e. reduce number of pixels and thereby resolution per pixel

Let us discuss each of these steps in more detail. Every raw image was processed, whereas the processed images of size 512×512 were saved in a folder named by its category. Note that every raw image resulted in approximately 100 – 150 processed images, so that we ended up with more than 10,000 images for each dataset.

Within each category of the processed images we labelled a selected portion of the images, by moving them into the folder with the respective label name. The folder structure we used is the following, where pointy brackets '`<parameter>`' indicate a parameter and 'etc' stands for the three label folders.

```
{raw-images-}usgs-<pixels>-res<resolution>m
  └── semi-desert
      ├── label-0
      └── label-1
```

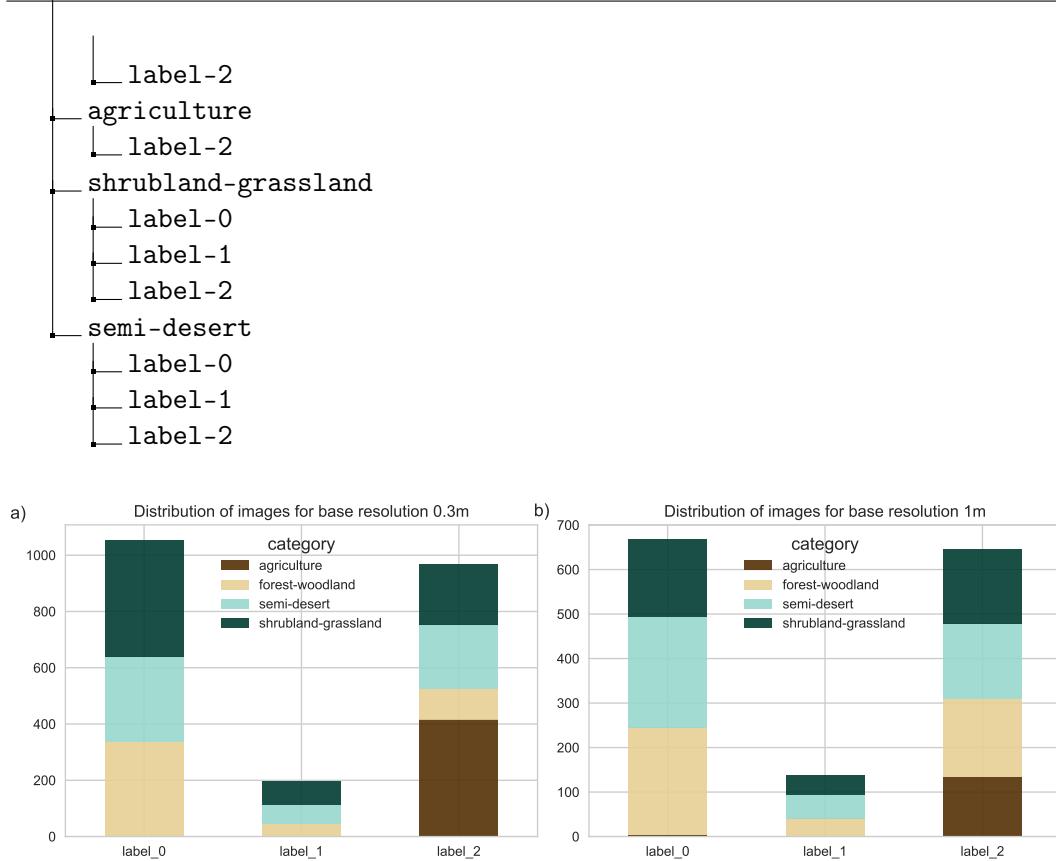


FIGURE 3.5: **Number of images per category and label.** (a) Distribution of images for dataset with resolution of 0.3m per pixel. (b) Distribution of images for dataset with resolution of 1m per pixel.

When labelling we stucked to the following rules. First, we classified images with no human impact at all into the class with label zero, while we classified images with very clear human influence into the class with label two. Ambigious images, i.e. images with minimal human trace were classified into label one. Second, we've put major effort into creating datasets that contain images of similar texture spread across all classes. If we for example classified a set of images of a certain forest type into the class with label zero we classified another set of images with a similar forest type, but containing a building or a street, into the class with label two. The same applies for images in class with label one when they contain e.g. a small walking path. We followed the latter rule for all categories except Agriculture. The Agriculture images all show human influence, and were therefore all classified with label two. By sticking to these rules, we are able to guarantee that the algorithm learns features that relate to human impact, and not to image artefacts such as color or texture.

In Figures 3.1 - 3.4 we display sample images for each of the four categories Agriculture, Shrubland-grassland, Forest-woodland and Semi-desert. These images belong to the dataset which has a resolution per pixel of 0.3m. Note that in Figs. 3.2 - 3.4 the first row represents images of label zero and the second row shows images that belong to label two. As mentioned above, the images in Fig. 3.1 all contain human influence.

The distribution of categories and labels is shown in Fig. 3.5. Overall, for the 0.3m dataset we classified about 2200 images, and for the 1m dataset we classified about 1450 images. During classification our main goal consisted in creating a balanced dataset between labels zero and two. The minority of images, roughly 10% of all classified images were assigned to label 1. These images were used at random to

investigate the behaviour of the Machine Learning classifier, which is discussed in chapter [4 VERIFY](#).

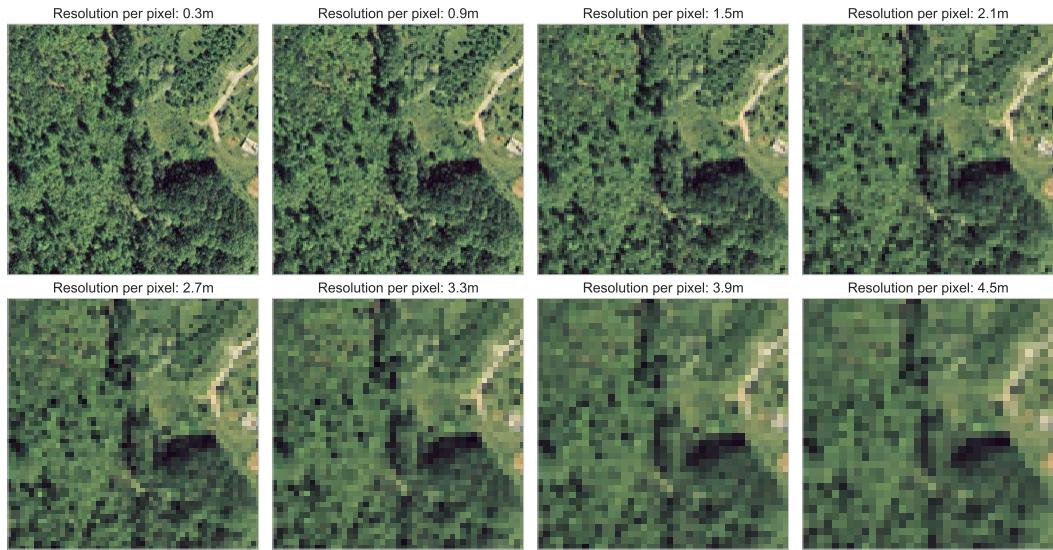


FIGURE 3.6: Example of image downsampling. The upper left image has the base resolution, 0.3m per pixel, and a size of 512×512 pixels whereas the lower right image has the worst resolution, 4.5m per pixel, and a size of 34×34 pixels. All intermediate images are downsampled by a factor corresponding to the resolution of the actual image divided by the base resolution. For instance, for the lower right image it is 15.

The last step of the data processing pipeline consisted in downsampling the processed and labelled images, to obtain images with a lower resolution per pixel. We used a Lanczos filter [6] for the sampling, which is based on a sinusoidal kernel. In Fig. 3.6 we show a few selected resolutions for an example image from the agriculture category. Note that here we only schematically depict an example in order to illustrate the process. However, in our Machine Learning pipeline the images are downsampled on the fly and the result of this process is not stored on disk.

For this particular image one can observe how certain image features disappear as the image quality is decreased. Above a resolution of around 3m per pixel one is not able anymore to identify the building close to the right corner of the image. The texture of the track that leads up to the building is blurred above a resolution of around 4m per pixel. This shows how different elements in an image are not recognizable anymore once the resolution is worse than their characteristic resolution.

Chapter 4

Deep Learning Approach

In the previous chapters we have introduced the key components of the approach we followed in our study: on the one hand, we have discussed Deep Learning and how it works, and on the other, we have described existing datasets and introduced the actual data we will consider. Now, we are ready to describe our approach, the image feature extraction, the model architecture and the training scenario.

4.1 Image features

In order to train a model based on images, features need to be extracted.

4.1.1 Traditional image features

Traditionally, image feature extraction was based on a set of hand-crafted detectors for edges (Sobel), corners, blobs (LoG, DoH) and other feature descriptors (SIFT, SURF, HOG)

- [https://en.wikipedia.org/wiki/Feature_detection_\(computer_vision\)](https://en.wikipedia.org/wiki/Feature_detection_(computer_vision))
- https://en.wikipedia.org/wiki/Scale-invariant_feature_transform
- https://en.wikipedia.org/wiki/Speeded_up_robust_features
- https://en.wikipedia.org/wiki/Histogram_of_oriented_gradients

(briefly describe, add proper references)

4.1.2 NN features and transfer learning

Recent approaches to image classification using Neural Networks have benefited from increasing computational power, and deep Convolutional Neural Networks have been able to achieve high performance.

Training a deep CNN from scratch requires a big and exhaustive dataset along with a huge amount of computational power, but it has been shown that the architectures of pre-trained NN can be reused for other purposes and achieve an equally great performance. Consequently, there is an increasing library of available pre-trained models and weights: Xception, VGG, InceptionV3, ResNet.

Most of these models have weights pre-trained on ImageNet, a large dataset containing more than 14 million hand-annotated images and over 20.000 categories.

These pre-trained architectures can be re-purposed by reusing the learnt weights and either replacing the final layers of the net by some other classifier, or even fine-tuning all the layers for the given problem. In any case, the initial layers of the Neural Network become a great image feature extractors.

- <https://iopscience.iop.org/article/10.1088/1742-6596/1087/6/062032/pdf>
- <https://arxiv.org/abs/1805.02294>
- <https://cs.nyu.edu/~fergus/papers/zeilerECCV2014.pdf>
- <https://towardsdatascience.com/transfer-learning-from-pre-trained-models-f2393f12>
- https://www.mitpressjournals.org/doi/pdf/10.1162/neco_a_00990
- <http://www.image-net.org/>
- <https://en.wikipedia.org/wiki/ImageNet>
- <https://keras.io/applications/>

4.1.3 pretrained resnet, weights on imagenet

architecture (image), training set, etc

- <https://towardsdatascience.com/review-resnet-winner-of-ilsvrc-2015-image-classification-5a2a2a2a>

4.2 Proposed architecture

4.2.1 resnet activations

images of activations (at several layers?)

-

4.2.2 complete architecture

architecture

4.2.3 training

cross validation

Chapter 5

Results

5.1 Results

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Aliquam ultricies lacinia euismod. Nam tempus risus in dolor rhoncus in interdum enim tincidunt. Donec vel nunc neque. In condimentum ullamcorper quam non consequat. Fusce sagittis tempor feugiat. Fusce magna erat, molestie eu convallis ut, tempus sed arcu. Quisque molestie, ante a tincidunt ullamcorper, sapien enim dignissim lacus, in semper nibh erat lobortis purus. Integer dapibus ligula ac risus convallis pellentesque.

5.1.1 Subsection 1

Nunc posuere quam at lectus tristique eu ultrices augue venenatis. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Aliquam erat volutpat. Vivamus sodales tortor eget quam adipiscing in vulputate ante ullamcorper. Sed eros ante, lacinia et sollicitudin et, aliquam sit amet augue. In hac habitasse platea dictumst.

5.2 Cost and Environment impact

Sed ullamcorper quam eu nisl interdum at interdum enim egestas. Aliquam placerat justo sed lectus lobortis ut porta nisl porttitor. Vestibulum mi dolor, lacinia molestie gravida at, tempus vitae ligula. Donec eget quam sapien, in viverra eros. Donec pellentesque justo a massa fringilla non vestibulum metus vestibulum. Vestibulum in orci quis felis tempor lacinia. Vivamus ornare ultrices facilisis. Ut hendrerit volutpat vulputate. Morbi condimentum venenatis augue, id porta ipsum vulputate in. Curabitur luctus tempus justo. Vestibulum risus lectus, adipiscing nec condimentum quis, condimentum nec nisl. Aliquam dictum sagittis velit sed iaculis. Morbi tristique augue sit amet nulla pulvinar id facilisis ligula mollis. Nam elit libero, tincidunt ut aliquam at, molestie in quam. Aenean rhoncus vehicula hendrerit.

Chapter 6

Conclusions

6.1 Conclusions

Lorem ipsum

6.1.1 Subsection 1

Nunc posuere

6.2 Further work

- Dataset: improve, enlarge, better classify, variety
- Model: CNN feature extraction (other architecture, activations)

Chapter 7

Introduction

7.0.1 Tables

Tables are an important way of displaying your results, below is an example table which was generated with this code:

```
\begin{table}
\caption{The effects of treatments X and Y on the four groups studied.}
\label{tab:treatments}
\centering
\begin{tabular}{l l l}
\toprule
\multicolumn{1}{c}{Groups} & \multicolumn{1}{c}{Treatment X} & \multicolumn{1}{c}{Treatment Y} \\
\midrule
1 & 0.2 & 0.8 \\
2 & 0.17 & 0.7 \\
3 & 0.24 & 0.75 \\
4 & 0.68 & 0.3 \\
\bottomrule
\end{tabular}
\end{table}
```

You can reference tables with `\ref{<label>}` where the label is defined within the table environment. See `Chapter1.tex` for an example of the label and citation (e.g. Table 7.1).

7.0.2 Figures

There will hopefully be many figures in your thesis (that should be placed in the *Figures* folder). The way to insert figures into your thesis is to use a code template like this:

```
\begin{figure}
```

TABLE 7.1: The effects of treatments X and Y on the four groups studied.

Groups	Treatment X	Treatment Y
1	0.2	0.8
2	0.17	0.7
3	0.24	0.75
4	0.68	0.3

```
\centering
\includegraphics{Figures/Electron}
\decoRule
\caption[An Electron]{An electron (artist's impression).}
\label{fig:Electron}
\end{figure}
```

Also look in the source file. Putting this code into the source file produces the picture of the electron that you can see in the figure below.

FIGURE 7.1: An electron (artist's impression).

Sometimes figures don't always appear where you write them in the source. The placement depends on how much space there is on the page for the figure. Sometimes there is not enough room to fit a figure directly where it should go (in relation to the text) and so L^AT_EX puts it at the top of the next page. Positioning figures is the job of L^AT_EX and so you should only worry about making them look good!

Figures usually should have captions just in case you need to refer to them (such as in Figure 7.1). The \caption command contains two parts, the first part, inside the square brackets is the title that will appear in the *List of Figures*, and so should be short. The second part in the curly brackets should contain the longer and more descriptive caption text.

The \decoRule command is optional and simply puts an aesthetic horizontal line below the image. If you do this for one image, do it for all of them.

L^AT_EX is capable of using images in pdf, jpg and png format.

7.0.3 Typesetting mathematics

If your thesis is going to contain heavy mathematical content, be sure that L^AT_EX will make it look beautiful, even though it won't be able to solve the equations for you.

The "Not So Short Introduction to L^AT_EX" (available on CTAN) should tell you everything you need to know for most cases of typesetting mathematics. If you need more information, a much more thorough mathematical guide is available from the AMS called, "A Short Math Guide to L^AT_EX" and can be downloaded from: <ftp://ftp.ams.org/pub/tex/doc/amsmath/short-math-guide.pdf>

There are many different L^AT_EX symbols to remember, luckily you can find the most common symbols in [The Comprehensive L^AT_EX Symbol List](#).

You can write an equation, which is automatically given an equation number by L^AT_EX like this:

```
\begin{equation}
E = mc^2
\label{eqn:Einstein}
\end{equation}
```

This will produce Einstein's famous energy-matter equivalence equation:

$$E = mc^2 \tag{7.1}$$

All equations you write (which are not in the middle of paragraph text) are automatically given equation numbers by L^AT_EX. If you don't want a particular equation numbered, use the unnumbered form:

```
\[ a^2=4 \]
```

7.1 Sectioning and Subsectioning

You should break your thesis up into nice, bite-sized sections and subsections. L^AT_EX automatically builds a table of Contents by looking at all the \chapter{}, \section{} and \subsection{} commands you write in the source.

The Table of Contents should only list the sections to three (3) levels. A chapter{} is level zero (0). A \section{} is level one (1) and so a \subsection{} is level two (2). In your thesis it is likely that you will even use a \subsubsection{}, which is level three (3). The depth to which the Table of Contents is formatted is set within `MastersDoctoralThesis.cls`. If you need this changed, you can do it in `main.tex`.

Bibliography

- [1] Saikat Basu et al. "DeepSat - A Learning framework for Satellite Imagery". In: *CoRR* abs/1509.03602 (2015). arXiv: [1509 . 03602](https://arxiv.org/abs/1509.03602). URL: <http://arxiv.org/abs/1509.03602>.
- [2] Leon Bottou and Olivier Bousquet. "The Tradeoffs of Large Scale Learning". In: (2008), pp. 161–168.
- [3] T Ciodaro et al. "Online particle detection with Neural Networks based on topological calorimetry information". In: *Journal of Physics: Conference Series* 368 (2012), p. 012030. DOI: [10 . 1088/1742-6596/368/1/012030](https://doi.org/10.1088/1742-6596/368/1/012030). URL: <https://doi.org/10.1088%2F1742-6596%2F368%2F1%2F012030>.
- [4] Ronan Collobert et al. "Natural Language Processing (almost) from Scratch". In: *CoRR* abs/1103.0398 (2011). arXiv: [1103 . 0398](https://arxiv.org/abs/1103.0398). URL: <http://arxiv.org/abs/1103.0398>.
- [5] G. Cybenko. "Approximation by superpositions of a sigmoidal function". In: *Mathematics of Control, Signals and Systems* 2.4 (1989), pp. 303–314. DOI: [10 . 1007/BF02551274](https://doi.org/10.1007/BF02551274). URL: <https://doi.org/10.1007/BF02551274>.
- [6] Claude E. Duchon. "Lanczos Filtering in One and Two Dimensions". In: *Journal of Applied Meteorology* 18, 1016-1022 (1979). URL: https://icess.eri.ucsb.edu/gem/Duchon_1979_JAM_Lanczos.pdf.
- [7] Xavier Glorot, Antoine Bordes, and Y Bengio. "Deep Sparse Rectifier Neural Networks". In: *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics (AISTATS) 2011* 15 (Jan. 2011), pp. 315–323.
- [8] *Google Maps: Conversion from zoom levels to distance.* <https://gis.stackexchange.com/questions/7430/what-ratio-scales-do-google-maps-zoom-levels-correspond-to>. accessed 17.06.2019.
- [9] Kaiming He et al. "Deep Residual Learning for Image Recognition". In: *CoRR* abs/1512.03385 (2015). arXiv: [1512 . 03385](https://arxiv.org/abs/1512.03385). URL: <http://arxiv.org/abs/1512.03385>.
- [10] Patrick Helber et al. "EuroSAT: A Novel Dataset and Deep Learning Benchmark for Land Use and Land Cover Classification". In: *CoRR* abs/1709.00029 (2017). arXiv: [1709 . 00029](https://arxiv.org/abs/1709.00029). URL: <http://arxiv.org/abs/1709.00029>.
- [11] G. Hinton et al. "Deep Neural Networks for Acoustic Modeling in Speech Recognition: The Shared Views of Four Research Groups". In: *IEEE Signal Processing Magazine* 29.6 (2012), pp. 82–97. ISSN: 1053-5888. DOI: [10 . 1109/MSP . 2012 . 2205597](https://doi.org/10.1109/MSP.2012.2205597).
- [12] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. "ImageNet Classification with Deep Convolutional Neural Networks". In: *Commun. ACM* 60 (2012), pp. 84–90.
- [13] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. In: *Nature* 521 (2015), 436 EP. URL: <https://doi.org/10.1038/nature14539>.

- [14] Junshui Ma et al. "Deep Neural Nets as a Method for Quantitative Structure: Activity Relationships". In: *Journal of Chemical Information and Modeling* 55.2 (2015). PMID: 25635324, pp. 263–274. DOI: [10.1021/ci500747n](https://doi.org/10.1021/ci500747n). eprint: <https://doi.org/10.1021/ci500747n>. URL: <https://doi.org/10.1021/ci500747n>.
- [15] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. "Learning representations by back-propagating errors". In: *Nature* 323.6088 (1986), pp. 533–536. DOI: [10.1038/323533a0](https://doi.org/10.1038/323533a0). URL: <https://doi.org/10.1038/323533a0>.
- [16] Gencer Sumbul et al. "BigEarthNet: A Large-Scale Benchmark Archive For Remote Sensing Image Understanding". In: *CoRR* abs/1902.06148 (2019). arXiv: [1902.06148](https://arxiv.org/abs/1902.06148). URL: [http://arxiv.org/abs/1902.06148](https://arxiv.org/abs/1902.06148).
- [17] Jonathan Tompson et al. "Joint Training of a Convolutional Network and a Graphical Model for Human Pose Estimation". In: *CoRR* abs/1406.2984 (2014). arXiv: [1406.2984](https://arxiv.org/abs/1406.2984). URL: [http://arxiv.org/abs/1406.2984](https://arxiv.org/abs/1406.2984).
- [18] Gui-Song Xia et al. "AID: A Benchmark Dataset for Performance Evaluation of Aerial Scene Classification". In: *CoRR* abs/1608.05167 (2016). arXiv: [1608.05167](https://arxiv.org/abs/1608.05167). URL: [http://arxiv.org/abs/1608.05167](https://arxiv.org/abs/1608.05167).
- [19] Yi Yang and Shawn Newsam. "Bag-of-visual-words and spatial extensions for land-use classification". In: Jan. 2010, pp. 270–279. DOI: [10.1145/1869790.1869829](https://doi.org/10.1145/1869790.1869829).
- [20] Weixun Zhou et al. "PatternNet: A Benchmark Dataset for Performance Evaluation of Remote Sensing Image Retrieval". In: *CoRR* abs/1706.03424 (2017). arXiv: [1706.03424](https://arxiv.org/abs/1706.03424). URL: [http://arxiv.org/abs/1706.03424](https://arxiv.org/abs/1706.03424).